

## 9-1. Logistic Regression Scratch

September 17, 2020

### 0.0.1 Algorithm 1: Logistic Regression

Logistic regression is a binary classification algorithm by simply finding a best fitted line that separates two dataset. In order to squash the output to a value between 0 and 1, logistic regression used a function called logit function (or sigmoid function)

Steps are the followings:

1. Prepare your data
  - add intercept
  - X and y and w in the right shape
    - X (m, n)
    - y (m, )
    - w (n, )
  - train-test split
  - feature scale
  - clean out any missing data
  - (optional) feature engineering
2. Predict and calculate the loss
  - The loss function is the cross entropy defined as

$$J = -\sum_{i=1}^m y^{(i)} \log(h) + (1 - y^{(i)}) \log(1 - h)$$

where h is defined as the sigmoid function as

$$h = \frac{1}{1 + e^{-\theta^T x}}$$

3. Calculate the gradient based on the loss
  - The gradient of  $\theta_j$  is defined as

$$\frac{\partial J}{\partial \theta_j} = \sum_{i=1}^m (h^{(i)} - y^{(i)}) x_j$$

- This can be derived by knowing that

$$J = y_1 \log h + (1 - y_1) \log(1 - h)$$

$$h = \frac{1}{1 + e^{-g}}$$
$$g = \theta^T x$$

- Thus, gradient of  $J$  in respect to some  $\theta_j$  is

$$\frac{\partial J}{\partial \theta_j} = \frac{\partial J}{\partial h} \frac{\partial h}{\partial g} \frac{\partial g}{\partial \theta_j}$$

where

$$\frac{\partial J}{\partial h} = \frac{y_1 - h}{h(1 - h)}$$

$$\frac{\partial h}{\partial g} = h(1 - h)$$

$$\frac{\partial g}{\partial \theta_j} = x_j$$

- Thus,

$$\begin{aligned} \frac{\partial J}{\partial \theta_j} &= \frac{\partial J}{\partial h} \frac{\partial h}{\partial g} \frac{\partial g}{\partial \theta_j} \\ &= \frac{y_1 - h}{h(1 - h)} * h(1 - h) * x_j \\ &= (y_1 - h)x_j \end{aligned}$$

- We can then put negative sign in front to make it negative loglikelihood, thus

$$(h - y_i)x_j$$

4. Update the theta with this update rule

$$\theta_j := \theta_j - \alpha * \frac{\partial J}{\partial \theta_j}$$

where  $\alpha$  is a typical learning rate range between 0 and 1

5. Loop 2-4 until max\_iter is reached, or the difference between old loss and new loss are smaller than some predefined threshold tol

## 0.0.2 Step 1: Prepare your data

### 1.1 Get your X and y in the right shape

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import linear_model


from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

#generate quite a lot of noise
#with only 4 informative features out of 10
#with 2 redundant features, overlapping with that 4 informative features
#and 4 noisy features
```

```
#Also, make std wider using n_clusters=2
X, y = make_classification(n_samples=500, n_features=10, n_redundant=2,
    ↪n_informative=4,
                                n_clusters_per_class=2, random_state=14)
plt.scatter(X[:, 0], X[:, 1], marker='o', c=y,
            s=25, edgecolor='k')
```

[1]: <matplotlib.collections.PathCollection at 0x12c71d610>



## 1.2 Feature scale your data to reach faster convergence

```
[2]: #feature scaling helps improve reach convergence faster
scaler = StandardScaler()
X = scaler.fit_transform(X)
```

## 1.3 Train test split your data

```
[3]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

## 1.4 Add intercepts

```
[4]: intercept = np.ones((X_train.shape[0], 1))
X_train = np.concatenate((intercept, X_train), axis=1) #add intercept
intercept = np.ones((X_test.shape[0], 1))
X_test = np.concatenate((intercept, X_test), axis=1) #add intercept
```

### 0.0.3 Step 2: Fit your algorithm

#### 1. Define your algorithm

```
[5]: # Your code here

import numpy as np
from sklearn.metrics import average_precision_score, classification_report

#here I use mini-batch as a demonstration
#you are free to use any variants of gradient descent
def mini_batch_GD(X, y, max_iter=1000):
    w = np.zeros(X.shape[1])
    l_rate = 0.01
    #10% of data
    batch_size = int(0.1 * X.shape[0])
    for i in range(max_iter):
        ix = np.random.randint(0, X.shape[0])
        batch_X = X[ix:ix+batch_size]
        batch_y = y[ix:ix+batch_size]
        w = w - l_rate * grad(batch_X, batch_y, w)
    return w, i

def grad(X, y, w):
    m = X.shape[0]
    h = h_theta(X, w)
    error = h - y
    return (1/m) * np.dot(X.T, error)

def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def h_theta(X, w):
    return sigmoid(X @ w)

def output(pred):
    return np.round(pred)

w, i = mini_batch_GD(X_train, y_train, max_iter=5000)
```

#### 2. Compute accuracy

```
[6]: yhat = output(h_theta(X_test, w))

print("APS: ", average_precision_score(y_test, yhat))
print("Report: ", classification_report(y_test, yhat))
```

APS: 0.7927592805204745

Report:	precision	recall	f1-score	support
---------	-----------	--------	----------	---------

	0	0.77	0.89	0.83	72
	1	0.88	0.76	0.81	78
accuracy				0.82	150
macro avg		0.83	0.82	0.82	150
weighted avg		0.83	0.82	0.82	150

#### 0.0.4 Algorithm 2: Multinomial Logistic Regression

This is logistic regression when number of classes are more than 2.

Steps are the followings:

The gradient descent has the following steps:

1. Prepare your data
  - add intercept
  - X and y in the right shape
    - X (m, n)
    - y (m, k)
    - w (n, k)
  - train-test split
  - feature scale
  - clean out any missing data
  - (optional) feature engineering
2. Predict and calculate the loss
  - The loss function is the cross entropy defined as

$$-\sum_{i=1}^m y^{(i)} \log(h)$$

where h is defined as the softmax function as

$$p(y = j \mid \theta^{(i)}) = \frac{e^{\theta^{(i)}}}{\sum_{k=1}^K e^{\theta_k^{(i)}}}$$

3. Calculate the gradient based on the loss
  - The gradient is defined as

$$\frac{\partial J}{\partial \theta_j} = \sum_{i=1}^m (h^{(i)} - y^{(i)}) x_j$$

- This gradient can be derived from the following simple examples:
  - Suppose given 2 classes (k = 2) and 3 features (n = 3), we have the loss function as

$$J = -y_1 \log h_1 - y_2 \log h_2$$

where  $h_1$  and  $h_2$  are

$$h_1 = \frac{\exp(g_1)}{\exp(g_1) + \exp(g_2)}$$

$$h_2 = \frac{\exp(g_2)}{\exp(g_1) + \exp(g_2)}$$

where  $g_1$  and  $g_2$  are

$$g_1 = w_{11}x_1 + w_{21}x_2 + w_{31}x_3$$

$$g_2 = w_{12}x_1 + w_{22}x_2 + w_{32}x_3$$

- For example, to find the gradient of  $J$  in respect to  $w_{21}$ , we simply can use the chain rule (or backpropagation) to calculate like this:

$$\frac{\partial J}{\partial w_{21}} = \frac{\partial J}{\partial h_1} \frac{\partial h_1}{\partial g_1} \frac{\partial g_1}{\partial w_{21}} + \frac{\partial J}{\partial h_2} \frac{\partial h_2}{\partial g_1} \frac{\partial g_1}{\partial w_{21}}$$

- If we know each of them, it is easy, where

$$\frac{\partial J}{\partial h_1} = -\frac{y_1}{h_1}$$

$$\frac{\partial J}{\partial h_2} = -\frac{y_2}{h_2}$$

$$\frac{\partial h_1}{\partial g_1} = \frac{\exp(g_1)}{\exp(g_1) + \exp(g_2)} - \left( \frac{\exp(g_1)}{\exp(g_1) + \exp(g_2)} \right)^2 = h_1(1 - h_1)$$

$$\frac{\partial h_2}{\partial g_1} = \frac{-\exp(g_2)\exp(g_1)}{(\exp(g_1) + \exp(g_2))^2} = -h_2h_1$$

$$\frac{\partial g_1}{\partial w_{21}} = x_2$$

- For those who forgets how to do third and fourth, recall that the quotient rule

$$\left( \frac{f}{g} \right)' = \frac{f'g - fg'}{g^2}$$

- Putting everything together, we got

$$\begin{aligned} \frac{\partial J}{\partial w_{21}} &= \frac{\partial J}{\partial h_1} \frac{\partial h_1}{\partial g_1} \frac{\partial g_1}{\partial w_{21}} + \frac{\partial J}{\partial h_2} \frac{\partial h_2}{\partial g_1} \frac{\partial g_1}{\partial w_{21}} \\ &= -\frac{y_1}{h_1} * h_1(1 - h_1) * x_2 + -\frac{y_2}{h_2} * -h_2h_1 * x_2 \\ &= x_2(-y_1 + y_1h_1 + y_2h_1) \\ &= x_2(-y_1 + h_1(y_1 + y_2)) \\ &= x_2(h_1 - y_1) \end{aligned}$$

4. Update the theta with this update rule

$$\theta_j := \theta_j - \alpha * \frac{\partial J}{\partial \theta_j}$$

where  $\alpha$  is a typical learning rate range between 0 and 1

5. Loop 2-4 until `max_iter` is reached, or the difference between old loss and new loss are smaller than some predefined threshold `tol`

### 0.0.5 Step 2: Fit your algorithm

#### 1. Define your algorithm

```
[7]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

#Step 1: Prepare data

# import some data to play with
iris = datasets.load_iris()
X = iris.data[:, 2:] # we only take the first two features.
y = iris.target #now our y is three classes thus require multinomial

#feature scaling helps improve reach convergence faster
scaler = StandardScaler()
X = scaler.fit_transform(X)


#data split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)

#add intercept to our X
intercept = np.ones((X_train.shape[0], 1))
X_train = np.concatenate((intercept, X_train), axis=1) #add intercept
intercept = np.ones((X_test.shape[0], 1))
X_test = np.concatenate((intercept, X_test), axis=1) #add intercept

#make sure our y is in the shape of (n, k)
#we will convert our output vector in
#matrix where no of columns is equal to the no of classes.
#The values in the matrix will be 0 or 1. For instance the rows
#where we have output 2 the column 2 will contain 1 and rest all 0.
#in simple words, y will be of shape (m, k)
k = len(set(y)) #no. of class (can also use np.unique)
m = X_train.shape[0] #no. of samples
n = X_train.shape[1] #no. of features
y_train_encoded = np.zeros((m, k))
for each_class in range(k):
    cond = y_train==each_class
    y_train_encoded[np.where(cond), each_class] = 1
```

```
[8]: #Step 1.1 (optional): Visualize our data

#your code here
plt.figure(figsize=(6,6))
plt.scatter(X[:, 0], X[:, 1], label='class 0', c=y)
plt.show()
```



[9]: *#Step 2: Fit your data*

```
import numpy as np
from sklearn.metrics import average_precision_score, classification_report
from sklearn.preprocessing import StandardScaler

def logistic_regression_GD(X, y, k, n, max_iter=1000):
    """
    Inputs:
        X shape: (m, n)
        w shape: (n, k)
    """
    w = np.random.rand(n, k)
    l_rate = 0.01
    for i in range(max_iter):
        cost, grad = gradient(X, y, w)
        w = w - l_rate * grad
    return w, i
```



```

#for those who tend to feel overwhelmed with lots of code
#I recommend you to write each part of the code separately as function
#it helps!
def gradient(X, y, w):
    m = X.shape[0]
    h = h_theta(X, w)
    cost = np.sum(-y * np.log(h) - (1 - y) * np.log(1 - h)) / m
    error = h - y
    grad = softmax_grad(X, error)
    return cost, grad

def softmax(x):
    return np.exp(x) / np.sum(np.exp(x), axis=1, keepdims=True)

def softmax_grad(X, error):
    return X.T @ error

def h_theta(X, w):
    '''
    Input:
        X shape: (m, n)
        w shape: (n, k)
    Returns:
        yhat shape: (m, k)
    '''
    # print("X@w: ", (X @ w)[:5])
    # print("Softmax: ", softmax(X @ w)[:5])
    return softmax(X @ w)

w, i = logistic_regression_GD(X_train, y_train_encoded, k, X_train.shape[1],
    ↪max_iter=5000)

pred = np.argmax(h_theta(X_test, w), axis=1)

print("Report: ", classification_report(y_test, pred))

```

Report:	precision	recall	f1-score	support
0	1.00	1.00	1.00	16
1	1.00	0.87	0.93	15
2	0.88	1.00	0.93	14
accuracy			0.96	45
macro avg	0.96	0.96	0.95	45
weighted avg	0.96	0.96	0.96	45

[ ]: