

5 - Pandas

August 24, 2020

1 Programming for Data Science and Artificial Intelligence

1.1 5 Pandas

1.1.1 Readings:

- [VANDER] Ch3
- <https://pandas.pydata.org/docs/>

Pandas is a newer package **built on top of NumPy**, and provides an efficient implementation of a DataFrame.

DataFrames are essentially multidimensional arrays with attached row and column labels, and often with heterogeneous types and/or missing data

Offering a convenient storage interface for labeled data, Pandas implements a number of powerful data operations familiar to users of both database frameworks and spreadsheet programs

Although NumPy ndarray data structure provides essential features, its limitations become clear when we need more flexibility (e.g., attaching labels to data, working with missing data, etc.)

Pandas provides efficient access to these sorts of “data munging” tasks that occupy much of a data scientist’s time.

```
[1]: import pandas as pd
      pd.__version__
```

```
[1]: '1.0.4'
```

1.2 Series

A Pandas Series is a one-dimensional array of indexed data. It can be created from a list or array as follows:

```
[2]: #creating a series from list
      data_pd = pd.Series([0.25, 0.5, 0.75, 1.0])
      print("1: Pandas series from list: ", data_pd)

      #creating a series from numpy array
      import numpy as np
      numpy_arr = np.arange(5)
      data_pd = pd.Series(numpy_arr)
```

```

print("2: Pandas series from numpy array: ", data_pd) #notice default index

#attributes
print("3: Data values: ", data_pd.values) #numpy array
print("4: Data index: ", data_pd.index) #range index is pandas object

#indexing
print("5: Data[1]: ", data_pd[1])
print("6: Data[-2:]: ", data_pd[-2:]) #slicing still works

#index is explicitly defined! Unlike numpy which is implicitly defined
#errors, no index of -1
# print("Data[-1]: ", data_pd[-1]) #you can provide index to pd.Series for
↳ this to work

#integer indexing is nasty! try change to letters
data_pd = pd.Series([0.25, 0.5, 0.75, 1.0],
                    index=['a', 'b', 'c', 'd'])

print("7: Data[-1]: ", data_pd[-1]) #you can provide index to pd.Series for
↳ this to work

#for example, index need not be integer
index = ['a', 'b', 'c', 'd', 3]
data_pd = pd.Series(numpy_arr, index = index)
print("8: Index['a']: ", data_pd['a'])
print("9: Index[3]: ", data_pd[3]) #3 is not position three but index 3!

#since pandas index system looks like dictionary, it is no surprise
#that pandas can be created from dictionary, see this:
some_population_dict = {'Chiangrai': 11111,
                        'Pathum Thani': 22222,
                        'Bangkok': 33333,
                        'Chiangmai': 44444}
data_pd = pd.Series(some_population_dict)
print("10: Population['Bangkok']: ", data_pd['Bangkok'])

#also supports slicing! Notice that Chiangmai is being included as well!
print("11: Population['Pathumthan':'Chiangmai']: ", data_pd['Pathum Thani':
↳ 'Chiangmai'])

#data can be scalar, which is repeated to fill the specified index
data_pd = pd.Series(5, index=[2, 3, 8])
print("12: Data pd scalar: ", data_pd)

```

```
1: Pandas series from list: 0    0.25
```

```
1    0.50
```

```

2    0.75
3    1.00
dtype: float64
2: Pandas series from numpy array:  0    0
1    1
2    2
3    3
4    4
dtype: int64
3: Data values:  [0 1 2 3 4]
4: Data index:  RangeIndex(start=0, stop=5, step=1)
5: Data[1]:  1
6: Data[-2:]:  3    3
4    4
dtype: int64
7: Data[-1]:  1.0
8: Index['a']:  0
9: Index[3]:  4
10: Population['Bangkok']:  33333
11: Population['Pathumthan':'Chiangmai']:  Pathum Thani    22222
Bangkok          33333
Chiangmai        44444
dtype: int64
12: Data pd scalar:  2    5
3    5
8    5
dtype: int64

```

1.3 Dataframe

DataFrame is an analog of a two-dimensional array with both flexible row indices and flexible column names. DataFrame as a sequence of aligned Series objects.

```

[3]: #create dataframe with two dict
some_population_dict = {'Chiangrai': 11111,
                        'Pathum Thani': 22222,
                        'Bangkok': 33333,
                        'Chiangmai': 44444}
some_area_dict = {'Chiangrai': 999,
                  'Pathum Thani': 888,
                  'Bangkok': 777,
                  'Chiangmai': 666,
                  'Syria': 333}
states = pd.DataFrame({'population': some_population_dict,
                       'area': some_area_dict})
print("Everything: ")
print(states)
print()

```

```

print("Only Chiangrai to Bangkok: ")
print(states['Chiangrai': 'Bangkok'])
print()

print("Only population column: ")
print(states['Chiangrai': 'Bangkok']['population'])
print(states['population']['Chiangrai': 'Bangkok']) #order does not matter;
    ↳ imagine as drill down
# print(states['population']['Chiangrai': 'Bangkok']['area']) #error, since
    ↳ area is gone already
print()

print("Only area column with everything")
print(states['area']) #first way
print(states[:, 'area']) #second way
print()

#attributes
print("Index: ", states.index) #pandas index object
print("Index[-1]: ", states.index[-1]) #pandas index object is similar to
    ↳ numpy array
print("Columns: ", states.columns) #pandas index object
print("Columns[0:1]: ", states.columns[0:1]) #notice how 1 is not included
print()

#many other ways to create panda dataframe

#from series
population_series = pd.Series(some_population_dict)
pd_from_series = pd.DataFrame(population_series, columns=['population'])
    ↳ #since it's series, we need to pass column name, if not, it will be named 0
print("PD from series: ", pd_from_series)

#from list of dicts
data = [{'a': i, 'b': 2 * i} for i in range(3)]
pd_from_list_dict = pd.DataFrame(data, index=[1, 2, 3])
print("PD from list of dict: ", pd_from_list_dict)

#from 2D numpy array
data_numpy = np.random.rand(3, 2)
index = ['a', 'b', 'c']
columns = ['foo', 'bar']
pd_from_numpy = pd.DataFrame(data_numpy, index=index, columns=columns)
print("PD from numpy: ", pd_from_numpy)

```

Everything:

	population	area
Chiangrai	11111.0	999
Pathum Thani	22222.0	888
Bangkok	33333.0	777
Chiangmai	44444.0	666
Syria	NaN	333

Only Chiangrai to Bangkok:

	population	area
Chiangrai	11111.0	999
Pathum Thani	22222.0	888
Bangkok	33333.0	777

Only population column:

Chiangrai	11111.0
Pathum Thani	22222.0
Bangkok	33333.0

Name: population, dtype: float64

Chiangrai	11111.0
Pathum Thani	22222.0
Bangkok	33333.0

Name: population, dtype: float64

Only area column with everything

Chiangrai	999
Pathum Thani	888
Bangkok	777
Chiangmai	666
Syria	333

Name: area, dtype: int64

Chiangrai	999
Pathum Thani	888
Bangkok	777
Chiangmai	666
Syria	333

Name: area, dtype: int64

Index: Index(['Chiangrai', 'Pathum Thani', 'Bangkok', 'Chiangmai', 'Syria'], dtype='object')

Index[-1]: Syria

Columns: Index(['population', 'area'], dtype='object')

Columns[0:1]: Index(['population'], dtype='object')

PD from series: population

Chiangrai	11111
Pathum Thani	22222
Bangkok	33333

```

Chiangmai          44444
PD from list of dict:      a  b
1  0  0
2  1  2
3  2  4
PD from numpy:          foo      bar
a  0.025979  0.236864
b  0.108672  0.139471
c  0.242305  0.780604

```

1.4 Indexing/Slicing/Fancy Indexing

1.4.1 Series

```

[4]: import pandas as pd
data = pd.Series([0.25, 0.5, 0.75, 1.0],
                  index = ['a', 'b', 'c', 'd'])

print("Data: ", data)

#explicit index
print("Explicit index: ", data['b'])

#also support implicit index, since we define index
print("Implicit index: ", data[-1])

#extend series by assigning new index value
data['e'] = 99
print("Data with e: ", data)

#use in
print("a in data?: ", 'a' in data) #access keys

#keys()
print("All keys: ", data.keys()) #pandas index object

#items()
print("All items: ", list(data.items())) #zip object

#values()
print("All values: ", data.values) #numpy array object

#slicing with explicit index (include c)
print("Data['a':'c']: ", data['a':'c'])

#slicing with implicit index (exclude 2)
print("Data[0:2]: ", data[0:2])

```

```

#masking
print("Data[(data > 0.3) & (data < 0.8)]: ", data[(data>0.3) & (data < 0.8)])

#fancy indexing
print("Data[['a', 'e']]: ", data[['a', 'e']])

#some precaution on the possible confusion if you use explicit integer index
#use explicit index when indexing
data = pd.Series(['a', 'b', 'c'], index=(1, 3, 5))
print("Data[1]: ", data[1])

#use implicit index when slicing
print("Data[1:3]: ", data[1:3])

#due to this possible confusion due to explicit integer index, pandas
#provide loc and iloc
#loc uses explicit index
print("Data loc [1]: ", data.loc[1])
print("Data loc [1:3]: ", data.loc[1:3]) #include 3 as well

#iloc uses implicit index
print("Data iloc [1]: ", data.iloc[1])
print("Data iloc [1:3]: ", data.iloc[1:3]) #does not include 3

#always use iloc for implicit index, for no unnecessary confusion!

```

```

Data:  a    0.25
      b    0.50
      c    0.75
      d    1.00
dtype: float64
Explicit index:  0.5
Implicit index:  1.0
Data with e:  a    0.25
              b    0.50
              c    0.75
              d    1.00
              e   99.00
dtype: float64
a in data?:  True
All keys:  Index(['a', 'b', 'c', 'd', 'e'], dtype='object')
All items:  [('a', 0.25), ('b', 0.5), ('c', 0.75), ('d', 1.0), ('e', 99.0)]
All values:  [ 0.25  0.5  0.75  1.  99. ]
Data['a':'c']:  a    0.25
              b    0.50
              c    0.75
dtype: float64

```

```

Data[0:2]:  a    0.25
           b    0.50
dtype: float64
Data[(data > 0.3) & (data < 0.8)]:  b    0.50
           c    0.75
dtype: float64
Data[['a', 'e']]:  a    0.25
                   e   99.00
dtype: float64
Data[1]:  a
Data[1:3]:  3    b
           5    c
dtype: object
Data loc [1]:  a
Data loc [1:3]:  1    a
                3    b
dtype: object
Data iloc [1]:  b
Data iloc [1:3]:  3    b
                 5    c
dtype: object

```

1.4.2 DataFrame

```

[5]: some_population_dict = {'Chiangrai': 11111,
                             'Pathum Thani': 22222,
                             'Bangkok': 33333,
                             'Chiangmai': 44444}
some_area_dict = {'Chiangrai': 999,
                  'Pathum Thani': 888,
                  'Bangkok': 777,
                  'Chiangmai': 666,
                  'Syria': 333}
data = pd.DataFrame({'pop': some_population_dict,
                     'area': some_area_dict})

print("==data==")
print(data)

#dataframe can be accessed via dictionary style indexing
print("==Area==")
print(data['area'])

#we can also use attribute-style
print("==Area using attributes==")
print(data.area)
print("Data area is same: ", data.area is data['area'])

```



```

#However!!, keep in mind that attribute style does not always work
#for example, pop is a method, thus data.pop will point to some method instead
print("==Do not use .pop!==" )
print("Some pop method: ", data.pop)
print("Data pop not the same: ", data.pop is data['pop'])

#reminder, index is columns, not index!
# print("Data['Chiangrai'] error: ", data['Chiangrai'])

#best practice is to use ['key'] style!

#feature engineer easily
print("==Feature engineering with density==")
data['density'] = data['pop'] / data['area']
print(data)

#we cannot use data[0] since we do not have explicit index called 0, use iloc
↳ instead
print("==First row using iloc==")
print(data.iloc[0])

print("==First three rows, first two columns using iloc==")
print(data.iloc[:3, :2])

#use loc for explicit index
print("==Use loc for explicit index")
print(data.loc[:'Bangkok', :'pop'])

#chaining is possible for combine
print("==chain iloc and loc==")
print(data.loc[:'Bangkok'].iloc[:, :2])

#more advanced: masking + fancy
print("==masking + fancy==")
print(data.loc[data.density > 20, ['pop', 'density']])

#first row, second column, change value to 90
data.iloc[0, 2] = 90
print("==Changing value==")
print(data)

###some really important convention to keep in mind###
#1. indexing refers to columns, while slicing refer to rows
print("==Indexing refer to columns==")
print(data['pop'])
#print(data['Chiangrai']) error

```

```

print("==Slicing refer to rows==")
print(data['Chiangrai': 'Pathum Thani']) #include Pathum Thani
#print(data['pop': 'area']) error

print("==Slicing can also use implicit integers==")
print(data[1:3]) #not including 3

print("==masking are done row-wise==")
print(data[data.density > 20])

```

==data==

	pop	area
Chiangrai	11111.0	999
Pathum Thani	22222.0	888
Bangkok	33333.0	777
Chiangmai	44444.0	666
Syria	NaN	333

==Area==

Chiangrai	999
Pathum Thani	888
Bangkok	777
Chiangmai	666
Syria	333

Name: area, dtype: int64

==Area using attributes==

Chiangrai	999
Pathum Thani	888
Bangkok	777
Chiangmai	666
Syria	333

Name: area, dtype: int64

Data area is same: True

==Do not use .pop!==

Some pop method: <bound method NDFrame.pop of

pop area

Chiangrai	11111.0	999
Pathum Thani	22222.0	888
Bangkok	33333.0	777
Chiangmai	44444.0	666
Syria	NaN	333>

Data pop not the same: False

==Feature engineering with density==

	pop	area	density
Chiangrai	11111.0	999	11.122122
Pathum Thani	22222.0	888	25.024775
Bangkok	33333.0	777	42.899614
Chiangmai	44444.0	666	66.732733

```

Syria          NaN    333          NaN
==First row using iloc==
pop           11111.000000
area          999.000000
density       11.122122
Name: Chiangrai, dtype: float64
==First three rows, first two columns using iloc==
      pop  area
Chiangrai  11111.0    999
Pathum Thani  22222.0    888
Bangkok      33333.0    777
==Use loc for explicit index
      pop
Chiangrai  11111.0
Pathum Thani  22222.0
Bangkok      33333.0
==chain iloc and loc==
      pop  area
Chiangrai  11111.0    999
Pathum Thani  22222.0    888
Bangkok      33333.0    777
==masking + fancy==
      pop  density
Pathum Thani  22222.0  25.024775
Bangkok      33333.0  42.899614
Chiangmai    44444.0  66.732733
==Changing value==
      pop  area  density
Chiangrai  11111.0    999  90.000000
Pathum Thani  22222.0    888  25.024775
Bangkok      33333.0    777  42.899614
Chiangmai    44444.0    666  66.732733
Syria        NaN    333          NaN
==Indexing refer to columns==
Chiangrai    11111.0
Pathum Thani  22222.0
Bangkok      33333.0
Chiangmai    44444.0
Syria        NaN
Name: pop, dtype: float64
==Slicing refer to rows==
      pop  area  density
Chiangrai  11111.0    999  90.000000
Pathum Thani  22222.0    888  25.024775
==Slicing can also use implicit integers==
      pop  area  density
Pathum Thani  22222.0    888  25.024775
Bangkok      33333.0    777  42.899614

```

==masking are done row-wise==

	pop	area	density
Chiangrai	11111.0	999	90.000000
Pathum Thani	22222.0	888	25.024775
Bangkok	33333.0	777	42.899614
Chiangmai	44444.0	666	66.732733

1.5 Broadcasting

```
[6]: import pandas as pd
import numpy as np

rng = np.random.RandomState(42)
df = pd.DataFrame(rng.randint(0, 10, (3, 4)),
                  columns = ['A', 'B', 'C', 'D'])
df
```

```
[6]:    A  B  C  D
0   6  3  7  4
1   6  9  2  6
2   7  4  3  7
```

```
[7]: #broadcasting
df_new = np.sin(df * np.pi / 4)
df_new

#index alignment
area = pd.Series({'Alaska': 111, 'Texas': 222,
                  'California': 333}, name = 'area')
population = pd.Series({'California': 999, 'Texas': 888,
                       'New York': 777}, name = 'population')

#as you can see, any missing value will be replaced with NaN
print("==population / area==")
print(population / area)

#similarly
A = pd.Series([2, 4, 6], index=[0, 1, 2])
B = pd.Series([1, 3, 5], index=[1, 2, 3])
print("==A + B==")
print(A + B)

#we can use fill_value params in pd.add(pd)
print("==A.add(B, fill_value=0)==") #any missing Nan values will be replaced
    ↳with 0
print(A.add(B, fill_value = 0))
```

```

#Operation between df and series
#common operation is to find difference of the entire df with one row or column
A = rng.randint(10, size = (3, 4))
df = pd.DataFrame(A, columns=list('QRST'))
print("==df==")
print(df)

#just simple question, why cannot we do df[0]
#print(df[0]) #by default, indexing access column wise

print("==df-df.iloc[0]==")
print(df - df.iloc[0])

print("==column-wise==") #cannot do -, since default is row-wise
print(df.subtract(df['R'], axis=0))

```

```

==population / area==
Alaska      NaN
California   3.0
New York     NaN
Texas        4.0
dtype: float64
==A + B==
0      NaN
1      5.0
2      9.0
3      NaN
dtype: float64
==A.add(B, fill_value=0)==
0      2.0
1      5.0
2      9.0
3      5.0
dtype: float64
==df==
   Q  R  S  T
0  7  2  5  4
1  1  7  5  1
2  4  0  9  5
==df-df.iloc[0]==
   Q  R  S  T
0  0  0  0  0
1 -6  5  0 -3
2 -3 -2  4  1
==column-wise==
   Q  R  S  T
0  5  0  3  2

```

```
1 -6  0 -2 -6
2  4  0  9  5
```

1.6 Handling missing data

```
[8]: #None
#Because it is a Python object, None cannot be used in any arbitrary
#NumPy/Pandas array, but only in arrays with data type 'object'
#(i.e., arrays of Python objects):

import numpy as np
import pandas as pd

vals = np.array([1, None, 3, 4])
vals

#typically, we do not use None, since it uses a lot of overhead
# for dtype in ['object', 'int']:
#     print("dtype=", dtype)
#     %timeit np.arange(1000, dtype=dtype).sum()

#in addition, you cannot perform aggregations like sum() or min() with
#array containing None
# vals.sum()

#instead, it is encouraged to use np.nan to represent null values
print("Type of np.nan: ", type(np.nan))
print("Np nan - 1: ", np.nan - 1) #any operation with nan is nan

vals2 = np.array([1, np.nan, 3, 4])
print("Sum: ", vals2.sum())
print("Nansum: ", np.nansum(vals2))
print("Nanmin: ", np.nanmin(vals2))
print("Nanmax: ", np.nanmax(vals2))

#Create a pd series. Automatically convert None to np.nan, if other values are
↳ integers or float
dfs = pd.Series([1, None, np.nan])
print("Pandas automatically convert None to nans")
print(dfs)

dfs = pd.Series(["Hello", None, np.nan])
print("Pandas do not, since String is object")
print(dfs)

dfs = pd.Series([True, None, np.nan])
print("Pandas do not, cast Boolean to object")
```

```

print(dfs)

#pandas treat None and np.nan interchangeably
#isnull() - check any missing values
#notnull() - opposite of isnull()
#dropna() - drop all null and return the filtered version
#fillna() - fill all null with some values
dfs = pd.Series([1, np.nan, "hello", None])
print("Is null: ", dfs.isnull())

print("==Only data not null==")
print(dfs[dfs.notnull()])

print("==Drop all na==")
print(dfs.dropna()) #this is not in place!
print(dfs) #this will bring back the old copy

#drop columns with na, use axis = 1
df = pd.DataFrame([[1, np.nan, 2],
                  [2, 3, 5],
                  [np.nan, 4, np.nan],
                  [4, np.nan, np.nan]])

print(df.dropna(axis=1)) #default is axis = 0

#drop columns, if all values is nan
df[3] = np.nan #create new column
print(df.dropna(axis=1, how="all"))

#most of the time, if you have missing values, it is advised
#to drop them simply. This is against many opinions. However,
#the case is that:
#When you replace missing value with mean, you lower the variance
#When you replace with median, you also lower the variance
#when you replace with some value, you introduce noise
#but first, let's see how to replace value
print("==Fill na with 0==")
print(df.fillna(0)) #not in place, reminders!

print("==Fill all na with mean==")
print(df.fillna(df.mean()))

print("replace df.mean() for col 1 with in place")
df[1].fillna(df[1].mean(), inplace=True) #this is in place!
print(df)

print("==Interpolate==")

```

```
print(df.interpolate(method='values'))
```

```
Type of np.nan: <class 'float'>
Np nan - 1: nan
Sum: nan
Nansum: 8.0
Nanmin: 1.0
Nanmax: 4.0
Pandas automatically convert None to nans
0    1.0
1    NaN
2    NaN
dtype: float64
Pandas do not, since String is object
0    Hello
1    None
2    NaN
dtype: object
Pandas do not, cast Boolean to object
0    True
1    None
2    NaN
dtype: object
Is null: 0    False
1    True
2    False
3    True
dtype: bool
==Only data not null==
0    1
2    hello
dtype: object
==Drop all na==
0    1
2    hello
dtype: object
0    1
1    NaN
2    hello
3    None
dtype: object
Empty DataFrame
Columns: []
Index: [0, 1, 2, 3]
      0    1    2
0  1.0  NaN  2.0
1  2.0  3.0  5.0
```



```

2 NaN 4.0 NaN
3 4.0 NaN NaN
==Fill na with 0==
      0      1      2      3
0 1.0 0.0 2.0 0.0
1 2.0 3.0 5.0 0.0
2 0.0 4.0 0.0 0.0
3 4.0 0.0 0.0 0.0
==Fill all na with mean==
      0      1      2      3
0 1.000000 3.5 2.0 NaN
1 2.000000 3.0 5.0 NaN
2 2.333333 4.0 3.5 NaN
3 4.000000 3.5 3.5 NaN
replace df.mean() for col 1 with in place
      0      1      2      3
0 1.0 3.5 2.0 NaN
1 2.0 3.0 5.0 NaN
2 NaN 4.0 NaN NaN
3 4.0 3.5 NaN NaN
==Interpolate==
      0      1      2      3
0 1.0 3.5 2.0 NaN
1 2.0 3.0 5.0 NaN
2 3.0 4.0 5.0 NaN
3 4.0 3.5 5.0 NaN

```

1.7 Concatenating Datasets

```

[9]: #dataframe concatenation
data_numpy = np.random.rand(3, 2) #shape 3, 2
index = ['Bangkok', 'Chiangmai', 'Samut Prakan']
columns = ['Population', 'Area']
pd_from_numpy = pd.DataFrame(data_numpy, index=index, columns=columns)
print("==First dataframe==")
print(pd_from_numpy)

data_numpy2 = np.random.rand(4, 3)
index2 = ['Bangkok', 'Chiangmai', 'Samut Prakan', 'Pathum Thani']
columns2 = ['HDI', 'Temperature', 'GDP']
pd_from_numpy2 = pd.DataFrame(data_numpy2, index=index2, columns=columns2)
print("==Second dataframe==")
print(pd_from_numpy2)

#do you see something weird, the index is duplicated!
#this is because pandas preserve indices!
#however, the columns are not duplicated...as you can see

```

```

print("==Normal concat along axis 1==")
print(pd.concat([pd_from_numpy, pd_from_numpy2], axis=1))

#to remove Nan from good, we use inner join (only perserve intersected elements)
print("==Concat with join inner==")
print(pd.concat([pd_from_numpy, pd_from_numpy2], axis=1, join='inner'))

#let say that you have more information of other countries
data_numpy3 = np.random.rand(3, 5) #shape 3, 2
index3 = ['Chiangrai', 'Korat', 'KhonKhaen']
columns3 = ['Population', 'Area', 'HDI', 'GDP', 'Temperature']
pd_from_numpy3 = pd.DataFrame(data_numpy3, index=index3, columns=columns3)

print("==Concat along axis 1 and 0==")
df = pd.concat([pd_from_numpy, pd_from_numpy2], axis=1)
df2 = pd.concat([df, pd_from_numpy3])
print(df2)

```

==First dataframe==

	Population	Area
Bangkok	0.387120	0.365903
Chiangmai	0.010938	0.201330
Samut Prakan	0.348705	0.504628

==Second dataframe==

	HDI	Temperature	GDP
Bangkok	0.058668	0.189302	0.397452
Chiangmai	0.975407	0.304613	0.731767
Samut Prakan	0.000055	0.979668	0.629784
Pathum Thani	0.108813	0.755269	0.710328

==Normal concat along axis 1==

	Population	Area	HDI	Temperature	GDP
Bangkok	0.387120	0.365903	0.058668	0.189302	0.397452
Chiangmai	0.010938	0.201330	0.975407	0.304613	0.731767
Samut Prakan	0.348705	0.504628	0.000055	0.979668	0.629784
Pathum Thani	NaN	NaN	0.108813	0.755269	0.710328

==Concat with join inner==

	Population	Area	HDI	Temperature	GDP
Bangkok	0.387120	0.365903	0.058668	0.189302	0.397452
Chiangmai	0.010938	0.201330	0.975407	0.304613	0.731767
Samut Prakan	0.348705	0.504628	0.000055	0.979668	0.629784

==Concat along axis 1 and 0==

	Population	Area	HDI	Temperature	GDP
Bangkok	0.387120	0.365903	0.058668	0.189302	0.397452
Chiangmai	0.010938	0.201330	0.975407	0.304613	0.731767
Samut Prakan	0.348705	0.504628	0.000055	0.979668	0.629784
Pathum Thani	NaN	NaN	0.108813	0.755269	0.710328
Chiangrai	0.395005	0.698924	0.552497	0.332663	0.212243

Korat	0.746205	0.520743	0.260659	0.290187	0.616227
KhonKhaen	0.645887	0.736909	0.024196	0.727968	0.041306

1.8 Merging Datasets with ID

```
[10]: left = pd.DataFrame({'ID': ['001', '002', '003', '005'],
                           'DS': ['B', 'B', 'B', 'C+'],
                           'SAD': ['A', 'B', 'C+', 'F']})
left
```

```
[10]:      ID  DS SAD
0  001   B   A
1  002   B   B
2  003   B  C+
3  005  C+   F
```

```
[11]: right = pd.DataFrame({'ID': ['001', '002', '003', '004'],
                             'HCI': ['B+', 'A', 'A', 'B+'],
                             'SDQI': ['A', 'A', 'B+', 'B']})
right
```

```
[11]:      ID HCI SDQI
0  001  B+   A
1  002   A   A
2  003   A  B+
3  004  B+   B
```

```
[12]: #seems like 004 is gone! Very similar to inner join
result = pd.merge(left, right, on='ID') #default is how=inner
result
```

```
[12]:      ID DS SAD HCI SDQI
0  001   B   A  B+   A
1  002   B   B   A   A
2  003   B  C+   A  B+
```

```
[13]: #specify how=outer
result = pd.merge(left, right, on='ID', how="outer") #default is how=inner
result
```

```
[13]:      ID  DS  SAD  HCI  SDQI
0  001    B    A  B+    A
1  002    B    B    A    A
2  003    B  C+    A  B+
3  005  C+    F  NaN  NaN
4  004  NaN  NaN  B+    B
```

```
[14]: #specify how=left, or right
result = pd.merge(left, right, on='ID', how="left") #default is how=inner
result
```

```
[14]:      ID  DS SAD  HCI SDQI
0   001   B   A   B+   A
1   002   B   B   A   A
2   003   B  C+   A  B+
3   005  C+   F  NaN  NaN
```

1.9 Aggregation

```
[15]: df = pd.DataFrame([('bird', 'Falconiformes', 389.0),
                        ('bird', 'Psittaciformes', 24.0),
                        ('mammal', 'Carnivora', 80.2),
                        ('mammal', 'Primates', np.nan),
                        ('mammal', 'Carnivora', 58)],
                        index=['falcon', 'parrot', 'lion', 'monkey', 'leopard'],
                        columns=('class', 'order', 'max_speed'))

df
```

```
[15]:      class      order  max_speed
falcon   bird  Falconiformes    389.0
parrot   bird  Psittaciformes     24.0
lion     mammal    Carnivora     80.2
monkey   mammal    Primates      NaN
leopard   mammal    Carnivora     58.0
```

```
[16]: grouped = df.groupby('class') #return a special view as DataFrameGroupByObject
grouped.sum()
```

```
[16]:      max_speed
class
bird      413.0
mammal    138.2
```

```
[17]: grouped = df.groupby('order') #return a special view as DataFrameGroupByObject
grouped.sum()
```

```
[17]:      max_speed
order
Carnivora      138.2
Falconiformes  389.0
Primates        0.0
Psittaciformes  24.0
```

```
[18]: grouped = df.groupby(['class', 'order'])
grouped.sum()
```

```
[18]:
```

		max_speed
class	order	
bird	Falconiformes	389.0
	Psittaciformes	24.0
mammal	Carnivora	138.2
	Primates	0.0

```
[19]: df.groupby(['class'])['max_speed'].median()
```

```
[19]: class
bird      206.5
mammal     69.1
Name: max_speed, dtype: float64
```

```
[20]: import seaborn as sns
planets = sns.load_dataset('planets')
print("Shape: ", planets.shape)
print("First five rows: ")
print(planets.head())

print("==Built in Pandas simple aggregations==")
print(planets.mean())
# print(planets.median())
# print(planets.min())
# print(planets.max())
# print(planets.std())
# print(planets.var())
# print(planets.sum())

print("==Mean (axis=1)==")
print(planets.mean(axis=1))

print("==Describe==")
print(planets.describe())
```

Shape: (1035, 6)

First five rows:

	method	number	orbital_period	mass	distance	year
0	Radial Velocity	1	269.300	7.10	77.40	2006
1	Radial Velocity	1	874.774	2.21	56.95	2008
2	Radial Velocity	1	763.000	2.60	19.84	2011
3	Radial Velocity	1	326.030	19.40	110.62	2007
4	Radial Velocity	1	516.220	10.50	119.47	2009

==Built in Pandas simple aggregations==

```

number          1.785507
orbital_period   2002.917596
mass             2.638161
distance         264.069282
year             2009.070531

```

```
dtype: float64
```

```
==Mean (axis=1)==
```

```

0      472.160000
1      588.586800
2      559.488000
3      492.810000
4      531.238000

```

```
...
```

```

1030    545.735377
1031    539.653966
1032    546.297881
1033    576.531271
1034    568.296939

```

```
Length: 1035, dtype: float64
```

```
==Describe==
```

	number	orbital_period	mass	distance	year
count	1035.000000	992.000000	513.000000	808.000000	1035.000000
mean	1.785507	2002.917596	2.638161	264.069282	2009.070531
std	1.240976	26014.728304	3.818617	733.116493	3.972567
min	1.000000	0.090706	0.003600	1.350000	1989.000000
25%	1.000000	5.442540	0.229000	32.560000	2007.000000
50%	1.000000	39.979500	1.260000	55.250000	2010.000000
75%	2.000000	526.005000	3.040000	178.500000	2012.000000
max	7.000000	730000.000000	25.000000	8500.000000	2014.000000

```
[21]: planets.groupby('method').sum()
```

```

[21]:
          number  orbital_period      mass  distance \
method
Astrometry          2    1.262360e+03    0.00000    35.75
Eclipse Timing Variations    15    4.276480e+04   10.25000   1261.44
Imaging             50    1.418973e+06    0.00000   2166.91
Microlensing        27    2.207500e+04    0.00000  41440.00
Orbital Brightness Modulation    5    2.127920e+00    0.00000   2360.00
Pulsar Timing        11    3.671511e+04    0.00000   1200.00
Pulsation Timing Variations    1    1.170000e+03    0.00000     0.00
Radial Velocity      952    4.553151e+05  1341.65638  27348.11
Transit             776    8.377523e+03    1.47000  134242.77
Transit Timing Variations    9    2.393505e+02    0.00000   3313.00

```

```
year
```

```
method
```

Astrometry	4023
Eclipse Timing Variations	18090
Imaging	76347
Microlensing	46225
Orbital Brightness Modulation	6035
Pulsar Timing	9992
Pulsation Timing Variations	2007
Radial Velocity	1110158
Transit	798461
Transit Timing Variations	8050

```
[22]: planets.groupby('method')['orbital_period'].median()
```

```
[22]: method
Astrometry                631.180000
Eclipse Timing Variations 4343.500000
Imaging                   27500.000000
Microlensing              3300.000000
Orbital Brightness Modulation 0.342887
Pulsar Timing             66.541900
Pulsation Timing Variations 1170.000000
Radial Velocity           360.200000
Transit                   5.714932
Transit Timing Variations 57.011000
Name: orbital_period, dtype: float64
```

```
[23]: #The GroupBy object supports direct iteration over the groups,
#returning each group as a Series or DataFrame:
for (method, group) in planets.groupby('method'):
    print("{0:30s} shape={1}".format(method, group.shape))
```

Astrometry	shape=(2, 6)
Eclipse Timing Variations	shape=(9, 6)
Imaging	shape=(38, 6)
Microlensing	shape=(23, 6)
Orbital Brightness Modulation	shape=(3, 6)
Pulsar Timing	shape=(5, 6)
Pulsation Timing Variations	shape=(1, 6)
Radial Velocity	shape=(553, 6)
Transit	shape=(397, 6)
Transit Timing Variations	shape=(4, 6)

```
[24]: planets.describe()
```

```
[24]:
```

	number	orbital_period	mass	distance	year
count	1035.000000	992.000000	513.000000	808.000000	1035.000000
mean	1.785507	2002.917596	2.638161	264.069282	2009.070531

std	1.240976	26014.728304	3.818617	733.116493	3.972567
min	1.000000	0.090706	0.003600	1.350000	1989.000000
25%	1.000000	5.442540	0.229000	32.560000	2007.000000
50%	1.000000	39.979500	1.260000	55.250000	2010.000000
75%	2.000000	526.005000	3.040000	178.500000	2012.000000
max	7.000000	730000.000000	25.000000	8500.000000	2014.000000

GroupBy objects have `aggregate()`, `filter()`, `transform()`, and `apply()` methods that efficiently implement a variety of useful operations before combining the grouped data. For the purpose of the following subsections, we'll use this DataFrame:

```
[25]: #aggregate() can take a string, a function or a list thereof, and compute
#all aggregates at once
planets.groupby('method')['orbital_period'].aggregate(
    ['min', np.median, max, np.mean, np.std, 'count'])
```

```
[25]:
```

	min	median	max \
method			
Astrometry	246.360000	631.180000	1016.000000
Eclipse Timing Variations	1916.250000	4343.500000	10220.000000
Imaging	4639.150000	27500.000000	730000.000000
Microlensing	1825.000000	3300.000000	5100.000000
Orbital Brightness Modulation	0.240104	0.342887	1.544929
Pulsar Timing	0.090706	66.541900	36525.000000
Pulsation Timing Variations	1170.000000	1170.000000	1170.000000
Radial Velocity	0.736540	360.200000	17337.500000
Transit	0.355000	5.714932	331.600590
Transit Timing Variations	22.339500	57.011000	160.000000

	mean	std	count
method			
Astrometry	631.180000	544.217663	2
Eclipse Timing Variations	4751.644444	2499.130945	9
Imaging	118247.737500	213978.177277	12
Microlensing	3153.571429	1113.166333	7
Orbital Brightness Modulation	0.709307	0.725493	3
Pulsar Timing	7343.021201	16313.265573	5
Pulsation Timing Variations	1170.000000	NaN	1
Radial Velocity	823.354680	1454.926210	553
Transit	21.102073	46.185893	397
Transit Timing Variations	79.783500	71.599884	3

```
[26]: #take only elements that belong to group with x orbital period of std less than 3
↪3
new_planets = planets.groupby('method').filter(lambda x: x['orbital_period'].
↪std() < 3)
new_planets.head()
```



```
[26]:
```

	method	number	orbital_period	mass	distance	\
787	Orbital Brightness Modulation	2	0.240104	NaN	1180.0	
788	Orbital Brightness Modulation	2	0.342887	NaN	1180.0	
792	Orbital Brightness Modulation	1	1.544929	NaN	NaN	

	year
787	2011
788	2011
792	2013

```
[27]: #look at each record
new_planets.loc[787]
```

```
[27]: method          Orbital Brightness Modulation
number                    2
orbital_period          0.240104
mass                    NaN
distance                1180
year                    2011
Name: 787, dtype: object
```

```
[28]: #perform apply for each group
planets.groupby('method').apply(lambda x: x.describe())
```

```
[28]:
```

		number	orbital_period	mass	distance	\
method	Astrometry	count	2.00	2.000000	0.0	2.000000
		mean	1.00	631.180000	NaN	17.875000
		std	0.00	544.217663	NaN	4.094148
		min	1.00	246.360000	NaN	14.980000
		25%	1.00	438.770000	NaN	16.427500
	
Transit Timing Variations	min	2.00	22.339500	NaN	339.000000	
	25%	2.00	39.675250	NaN	597.000000	
	50%	2.00	57.011000	NaN	855.000000	
	75%	2.25	108.505500	NaN	1487.000000	
	max	3.00	160.000000	NaN	2119.000000	

	year		
method	Astrometry	count	2.000000
		mean	2011.500000
		std	2.12132
		min	2010.000000
		25%	2010.750000
	
Transit Timing Variations	min	2011.000000	
	

25%	2011.75000
50%	2012.50000
75%	2013.25000
max	2014.00000

[80 rows x 5 columns]

```
[29]: #we can also do something fancy like this
grouped= planets.groupby('method')['orbital_period']

def some_func(group):
    return pd.DataFrame({'original': group,
                        'demeaned': group - group.mean()})

grouped.apply(some_func)
```

```
[29]:      original    demeaned
0      269.300000 -554.054680
1      874.774000   51.419320
2      763.000000 -60.354680
3      326.030000 -497.324680
4      516.220000 -307.134680
...
1030     3.941507 -17.160566
1031     2.615864 -18.486209
1032     3.191524 -17.910549
1033     4.125083 -16.976990
1034     4.187757 -16.914316
```

[1035 rows x 2 columns]

1.10 Pivot Tables

```
[30]: titanic = sns.load_dataset('titanic')
titanic.head()
```

```
[30]:   survived  pclass    sex  age  sibsp  parch   fare embarked  class \
0         0        3   male  22.0     1     0   7.2500         S   Third
1         1        1  female  38.0     1     0  71.2833         C   First
2         1        3  female  26.0     0     0   7.9250         S   Third
3         1        1  female  35.0     1     0  53.1000         S   First
4         0        3   male  35.0     0     0   8.0500         S   Third

      who  adult_male  deck  embark_town  alive  alone
0   man         True  NaN  Southampton    no  False
1 woman        False    C   Cherbourg   yes  False
2 woman        False  NaN  Southampton   yes   True
```

3	woman	False	C	Southampton	yes	False
4	man	True	NaN	Southampton	no	True

To start learning more about this data, we might begin by grouping according to gender, survival status, or some combination thereof. If you have read the previous section, you might be tempted to apply a GroupBy operation—for example, let’s look at survival rate by gender:

```
[31]: titanic.groupby('sex')[['survived']].mean()
```

```
[31]:      survived
sex
female  0.742038
male    0.188908
```

we might like to go one step deeper and look at survival by both sex and, say, class. Using the vocabulary of GroupBy, we might proceed using something like this: we group by class and gender, select survival, apply a mean aggregate, combine the resulting groups. In code:

```
[32]: titanic.groupby(['sex', 'class'])['survived'].aggregate('mean')
```

```
[32]: sex      class
female First      0.968085
      Second     0.921053
      Third      0.500000
male   First      0.368852
      Second     0.157407
      Third      0.135447
Name: survived, dtype: float64
```

The code is getting messy and it’s not supposed to be like that. Let’s use pivot_table. Here is the equivalent to the preceding operation using the pivot_table method of DataFrames:

```
[33]: titanic.pivot_table('survived', index='sex', columns='class')
```

```
[33]: class      First      Second      Third
sex
female  0.968085  0.921053  0.500000
male    0.368852  0.157407  0.135447
```

Just as in the GroupBy, the grouping in pivot tables can be specified with multiple levels, and via a number of options. For example, we might be interested in looking at age as a third dimension. We’ll bin the age using the pd.cut function:

```
[34]: age = pd.cut(titanic['age'], [0, 18, 80]) #return tuples of bins
titanic.pivot_table('survived', ['sex', age], 'class')
```

```
[34]: class      First      Second      Third
sex      age
```

```
female (0, 18]    0.909091  1.000000  0.511628
        (18, 80]  0.972973  0.900000  0.423729
male   (0, 18]    0.800000  0.600000  0.215686
        (18, 80]  0.375000  0.071429  0.133663
```

We can apply the same strategy when working with the columns as well; let's add info on the fare paid using `pd.qcut` to automatically compute quantiles:

```
[35]: fare = pd.qcut(titanic['fare'], 2) #2 equal sized groupings of the data
      titanic.pivot_table('survived', ['sex', age], [fare, 'class'])
```

```
[35]: fare          (-0.001, 14.454]          (14.454, 512.329] \
      class                First      Second      Third      First
      sex    age
female (0, 18]          NaN  1.000000  0.714286          0.909091
        (18, 80]          NaN  0.880000  0.444444          0.972973
male   (0, 18]          NaN  0.000000  0.260870          0.800000
        (18, 80]          0.0  0.098039  0.125000          0.391304
```

```
fare
class          Second      Third
sex    age
female (0, 18]    1.000000  0.318182
        (18, 80]  0.914286  0.391304
male   (0, 18]    0.818182  0.178571
        (18, 80]  0.030303  0.192308
```

1.11 String operations

```
[36]: names = pd.Series(['Graham Chapman', ' John Cleese ', 'Terry Gilliam ',
                        'Eric Idle', 'Terry Jones', 'Michael Palin'])
```

```
names.str.len()
names.str.lower()
names.str.startswith('T')
names.str.split() #default is space
names.str.extract('([A-Za-z]+)')
names.str[0:3]
names.str.strip()
names.str.replace(r'man', 'something')
```

```
[36]: 0    Graham Chapsomething
      1         John Cleese
      2       Terry Gilliam
      3         Eric Idle
      4       Terry Jones
      5       Michael Palin
```

```
dtype: object
```

1.12 Time Series

Pandas was developed in the context of financial modeling, so as you might expect, it contains a fairly extensive set of tools for working with dates, times, and time-indexed data. Let's first understand how Python treat dates and times

```
[37]: from datetime import datetime
      datetime(year=2025, month=7, day=4)
```

```
[37]: datetime.datetime(2025, 7, 4, 0, 0)
```

Using the dateutil module, you can parse dates from a variety of string formats

```
[38]: from dateutil import parser
      date = parser.parse("4th of July, 2015")
      date
```

```
[38]: datetime.datetime(2015, 7, 4, 0, 0)
```

Once you have a datetime object, you can do things like printing the day of the week:

```
[39]: date.strftime('%A, %D')
```

```
[39]: 'Saturday, 07/04/15'
```

We can similarly create numpy using type `np.datetime64` which is a very efficient way to store datetime

```
[40]: date = np.array('2015-07-04', dtype=np.datetime64)  #ISO date
      date
```

```
[40]: array('2015-07-04', dtype='datetime64[D]')
```

Given its numpy type, we can quickly do vectorized operations on it

```
[41]: date + np.arange(12)
```

```
[41]: array(['2015-07-04', '2015-07-05', '2015-07-06', '2015-07-07',
          '2015-07-08', '2015-07-09', '2015-07-10', '2015-07-11',
          '2015-07-12', '2015-07-13', '2015-07-14', '2015-07-15'],
          dtype='datetime64[D]')
```

Pandas uses **Timestamp** object, which combines the ease-of-use of datetime and dateutil with the efficient storage and vectorized interface of `numpy.datetime64`. From a group of these Timestamp objects, Pandas can construct a **DatetimeIndex** that can be used to index data in a Series or DataFrame

```
[42]: import pandas as pd
date = pd.to_datetime("2015-07-04")
print(type(date))
date.strftime('%A')
```

```
<class 'pandas._libs.tslibs.timestamps.Timestamp'>
```

```
[42]: 'Saturday'
```

```
[43]: #perform numpy style vectorized operations using pd.to_timedelta
date + pd.to_timedelta(np.arange(12), 'D') #unit is nanoseconds #does not
↳ support Y and M because each M has unequal amount of nanoseconds
```

```
[43]: DatetimeIndex(['2015-07-04', '2015-07-05', '2015-07-06', '2015-07-07',
                    '2015-07-08', '2015-07-09', '2015-07-10', '2015-07-11',
                    '2015-07-12', '2015-07-13', '2015-07-14', '2015-07-15'],
                    dtype='datetime64[ns]', freq=None)
```

Where the Pandas time series tools really become useful is when you begin to index data by timestamps. For example, we can construct a Series object that has time indexed data:

```
[44]: # index = pd.DatetimeIndex(['2014-07-04', '2014-08-04',
#                               '2015-07-04', '2015-08-04'])
index = pd.to_datetime(['2014-07-04', '2014-08-04',
                        '2015-07-04', '2015-08-04'])
data = pd.DataFrame(np.random.rand(4,2), index = index, columns=['Apple',
↳ 'Orange'])
data
```

```
[44]:
```

	Apple	Orange
2014-07-04	0.676958	0.395209
2014-08-04	0.955737	0.233010
2015-07-04	0.422293	0.297442
2015-08-04	0.756082	0.993525

```
[45]: #use slicing to access rows
data['2014-07-04':'2015-07-04']
```

```
[45]:
```

	Apple	Orange
2014-07-04	0.676958	0.395209
2014-08-04	0.955737	0.233010
2015-07-04	0.422293	0.297442

```
[46]: #there are special date-only indexing, such as passing a year
data['2015']
```

```
[46]:
```

	Apple	Orange
2015-07-04	0.422293	0.297442
2015-08-04	0.756082	0.993525

A useful method is `date_range`, which generates date from specified start and end

```
[47]: pd.date_range('2015-07-03', '2015-07-10')
```

```
[47]: DatetimeIndex(['2015-07-03', '2015-07-04', '2015-07-05', '2015-07-06',
                    '2015-07-07', '2015-07-08', '2015-07-09', '2015-07-10'],
                    dtype='datetime64[ns]', freq='D')
```

```
[48]: pd.date_range('2015-07-03', periods = 8) #instead of end, we can specify the
      ↪ periods
```

```
[48]: DatetimeIndex(['2015-07-03', '2015-07-04', '2015-07-05', '2015-07-06',
                    '2015-07-07', '2015-07-08', '2015-07-09', '2015-07-10'],
                    dtype='datetime64[ns]', freq='D')
```

```
[49]: #if we want 8 periods, but in hours, we use freq params
      pd.date_range('2015-07-03', periods = 8, freq='H') #try M
```

```
[49]: DatetimeIndex(['2015-07-03 00:00:00', '2015-07-03 01:00:00',
                    '2015-07-03 02:00:00', '2015-07-03 03:00:00',
                    '2015-07-03 04:00:00', '2015-07-03 05:00:00',
                    '2015-07-03 06:00:00', '2015-07-03 07:00:00'],
                    dtype='datetime64[ns]', freq='H')
```

```
[50]: #if we want 8 periods, but in hours, we use freq params
      pd.date_range('2015-07-03', periods = 18, freq='MS') #BH - Business Hours, MS -
      ↪ Month start
```

```
[50]: DatetimeIndex(['2015-08-01', '2015-09-01', '2015-10-01', '2015-11-01',
                    '2015-12-01', '2016-01-01', '2016-02-01', '2016-03-01',
                    '2016-04-01', '2016-05-01', '2016-06-01', '2016-07-01',
                    '2016-08-01', '2016-09-01', '2016-10-01', '2016-11-01',
                    '2016-12-01', '2017-01-01'],
                    dtype='datetime64[ns]', freq='MS')
```

```
[51]: pd.date_range('2015-07-03', periods = 18, freq='2H30T') #supports custom numbers
```

```
[51]: DatetimeIndex(['2015-07-03 00:00:00', '2015-07-03 02:30:00',
                    '2015-07-03 05:00:00', '2015-07-03 07:30:00',
                    '2015-07-03 10:00:00', '2015-07-03 12:30:00',
                    '2015-07-03 15:00:00', '2015-07-03 17:30:00',
                    '2015-07-03 20:00:00', '2015-07-03 22:30:00',
                    '2015-07-04 01:00:00', '2015-07-04 03:30:00',
```

```

'2015-07-04 06:00:00', '2015-07-04 08:30:00',
'2015-07-04 11:00:00', '2015-07-04 13:30:00',
'2015-07-04 16:00:00', '2015-07-04 18:30:00'],
dtype='datetime64[ns]', freq='150T')

```

```

[52]: from pandas_datareader import data

goog = data.DataReader('GOOG', start='2004', end='2016',
                        data_source='yahoo')

goog.head()

```

```

/home/akrarads/.local/lib/python3.8/site-
packages/pandas_datareader/compat/__init__.py:7: FutureWarning:
pandas.util.testing is deprecated. Use the functions in the public API at
pandas.testing instead.
  from pandas.util.testing import assert_frame_equal

```

```

[52]:

```

	High	Low	Open	Close	Volume	Adj Close
Date						
2004-08-19	51.835709	47.800831	49.813286	49.982655	44871300.0	49.982655
2004-08-20	54.336334	50.062355	50.316402	53.952770	22942800.0	53.952770
2004-08-23	56.528118	54.321388	55.168217	54.495735	18342800.0	54.495735
2004-08-24	55.591629	51.591621	55.412300	52.239193	15319700.0	52.239193
2004-08-25	53.798351	51.746044	52.284027	52.802086	9232100.0	52.802086

```


[53]: import matplotlib.pyplot as plt
goog_close = goog['Close']
goog_close.plot()

```

```

[53]: <matplotlib.axes._subplots.AxesSubplot at 0x7f881cdebc70>

```


One common need for time series data is resampling at a higher or lower frequency. This can be done using the `resample()` method, or the much simpler `asfreq()` method. The primary difference between the two is that `resample()` is fundamentally a data aggregation, while `asfreq()` is fundamentally a data selection.

Taking a look at the Google closing price, let's compare what the two return when we down-sample the data. Here we will resample the data at the end of business year:

```
[54]: goog_close.plot(alpha=0.2, style='--')
      goog_close.resample('BA').mean().plot(style='o') #BA = Business Year End
      goog_close.asfreq('BA').plot(style='--')

      #Notice the difference: at each point, resample reports the average
      #of the previous year, while asfreq reports the value at the end of the
      #year.
```

```
[54]: <matplotlib.axes._subplots.AxesSubplot at 0x7f881d23a610>
```



```
[55]: #difference between resample and asfreq
ts = pd.Series(range(365), index = pd.date_range(start='20190101',
                                                end='20191231',
                                                freq = 'D'))

ts.head()
```

```
[55]: 2019-01-01    0
      2019-01-02    1
      2019-01-03    2
      2019-01-04    3
      2019-01-05    4
      Freq: D, dtype: int64
```

```
[56]: ts.asfreq(freq='Q') #Quarter
```

```
[56]: 2019-03-31    89
      2019-06-30   180
      2019-09-30   272
      2019-12-31   364
      Freq: Q-DEC, dtype: int64
```

```
[57]: #this is not the average of each quarter, but the average of the last
#day of each quarter
ts.asfreq(freq='Q').mean()  #(89+180+272+364)/4
```

```
[57]: 226.25
```


```
[58]: bins = ts.resample('Q') #return DateTimeIndexResampler which is like GroupBy, ↵  
      ↪ object  
      #it actually create some groups  
  
      #this is actually the average of each quarter  
      bins.groups  
      bins.mean() #think of resample as groupby, (you can call mean, sum, apply, ↵  
      ↪ just like group by)
```

```
[58]: 2019-03-31      44.5  
      2019-06-30     135.0  
      2019-09-30     226.5  
      2019-12-31     318.5  
      Freq: Q-DEC, dtype: float64
```

Pandas has two closely related methods: `shift()` and `tshift()`. In short, the difference between them is that `shift()` shifts the data, while `tshift()` shifts the index. In both cases, the shift is specified in multiples of the frequency. Here we will both `shift()` and `tshift()` by 900 days;

```
[59]: fig, ax = plt.subplots(3, sharey=True)  
  
      goog_close = goog_close.asfreq('D', method='pad')  
  
      goog_close.plot(ax=ax[0])  
      goog_close.shift(900).plot(ax=ax[1]) #shift the data  
      goog_close.tshift(900).plot(ax=ax[2]) #simply shift the index
```


```
[59]: <matplotlib.axes._subplots.AxesSubplot at 0x7f881d0fbe80>
```



```
[60]: #why shift?
#A common context for this type of shift is in computing differences
#over time. For example, we use shifted values to compute the one-year
#return on investment for Google stock over the course of the dataset:
ROI = 100 * (goog_close.tshift(-365) / goog_close)
ROI.plot()
plt.ylabel('Return on Investment')

'''
This helps us to see the overall trend in Google stock:
thus far, the most profitable times to invest in Google
have been (unsurprisingly, in retrospect) shortly after its IPO, and
in the middle of the 2009 recession.
'''
```

```
[60]: '\nThis helps us to see the overall trend in Google stock: \nthus far, the most
profitable times to invest in Google \nhave been (unsurprisingly, in retrospect)
shortly after its IPO, and \nin the middle of the 2009 recession.\n'
```




```
[61]: ## Rolling window
rolling = goog_close.rolling(365, center = True)
rolling
```

```
[61]: Rolling [window=365,center=True,axis=0]
```

```
[62]: data = pd.DataFrame({'input': goog_close,
                        'one_year rolling mean': rolling.mean(),
                        'one_year rolling std': rolling.std()})

data.plot(style=['-', '--', ':'])
```

```
[62]: <matplotlib.axes._subplots.AxesSubplot at 0x7f8818f8feb0>
```



```
[63]: #Example: Visualizing Seattle Bicycle Counts
data = pd.read_csv('resources/FremontBridge.csv')
data.head()

data = pd.read_csv('resources/FremontBridge.csv', index_col='Date',
                   parse_dates=True)
```

```
[64]: data.columns = ['West', 'East'] #rename
data.head()
```

```
[64]:
```

	West	East
Date		
2019-03-31 23:00:00	6.0	10.0
2019-03-31 22:00:00	7.0	14.0
2019-03-31 21:00:00	18.0	15.0
2019-03-31 20:00:00	26.0	31.0
2019-03-31 19:00:00	30.0	58.0

```
[65]: #check any missing values
data.isna().sum()

#fillna with 0
data.fillna(0, inplace=True)

#feature engineering
```

```
data['Total'] = data['East'] + data['West']
```

```
[66]: data.head()
```

```
[66]:
```

		West	East	Total
Date				
2019-03-31 23:00:00		6.0	10.0	16.0
2019-03-31 22:00:00		7.0	14.0	21.0
2019-03-31 21:00:00		18.0	15.0	33.0
2019-03-31 20:00:00		26.0	31.0	57.0
2019-03-31 19:00:00		30.0	58.0	88.0


```
[67]: #let do some plot, checking on weekly trend

import matplotlib.pyplot as plt
import numpy as np

weekly = data.resample('W').agg('sum') #can also do .sum()
weekly.plot(style=[":", "--", "-"])
plt.ylabel('Weekly bicycle count')


#seems like people bicycle in summer more than winter
```

```
[67]: Text(0, 0.5, 'Weekly bicycle count')
```




```
[68]: #let's look at what time people usually bicycle
by_time = data.groupby(data.index.time).mean()
by_time.plot()
# hourly_ticks = 4 * 60 * 60 * np.arange(6) #every four hours
# by_time.plot(xticks=hourly_ticks, style=[':', '--', '-'])
```

[68]: <matplotlib.axes._subplots.AxesSubplot at 0x7f8818e9cd00>



```
[69]: #how about days?
by_day = data.groupby(data.index.dayofweek).mean()
by_day.index = ['Mon', 'Tues', 'Wed', 'Thurs', 'Fri', 'Sat', 'Sun']
by_day.plot(style=[':', '--', '-'])
```

[69]: <matplotlib.axes._subplots.AxesSubplot at 0x7f8818ccfd60>



[]: